



Evaluation Test Plan

Request Tracker Agent

10-Question Single Response Test Set

Last Updated: April 15, 2026

Developer: Power CAT

Agent Type: Copilot Studio Custom Agent

Table of Contents

Table of Contents	2
 1. Overview & Agent Summary	3
Agent Purpose	3
Knowledge Sources & Capabilities	3
Conversation Starters	4
Special Modes	4
Key AI Settings	4
 2. Test Objectives	4
 3. Test Cases	5
Test Case 1:  Core Knowledge	5
Test Case 2:  Core Knowledge	5
Test Case 3:  Conversation Starters	6
Test Case 4:  Action / Routing	6
Test Case 5:  Action / Routing	7
Test Case 6:  Specificity & Precision	7
Test Case 7:  Edge Cases & Boundaries	8
Test Case 8:  Multi-Part / Complex	8
Test Case 9:  Rephrased / Paraphrased	8
Test Case 10:  Safety & Guardrails	9
 4. Test Coverage & Method Reference	9
Category Distribution	9
Test Method Distribution	10
Coverage Notes	10
Test Method Reference	10
 5. Tips for Running Evaluations	11
How to Use the Evaluations Tab in Copilot Studio	11
Before Running	12
During and After	12
Iterating on the Test Set	12
 6. Resources	12

1. Overview & Agent Summary

This document provides a comprehensive evaluation test plan for the Request Tracker Agent — a custom agent built in Microsoft Copilot Studio. The test set contains 10 carefully crafted questions designed to validate the agent's core capabilities: request creation, status updates, search & query, topic routing accuracy, pronoun resolution, ambiguity handling, and safety guardrails. A companion CSV file is provided for direct import into Copilot Studio's Evaluation tab.

Field	Value
Agent Name	Request Tracker Agent
Agent Type	Copilot Studio Custom Agent
Developer	Power CAT
Test Set Size	10 Questions
Test Set Type	Single Response Evaluation
Last Updated	April 15, 2026
Companion File	Request Tracker Agent Evaluation Test Set.csv

Agent Purpose

An intelligent, cross-department request management assistant that enables employees to submit, track, escalate, and resolve requests through natural conversation. The agent proactively searches for existing resolutions before creating new tickets, auto-categorizes requests using AI, manages approval workflows via Power Automate, and keeps stakeholders informed with real-time notifications and scheduled digests. The agent replaces fragmented request intake processes — shared mailboxes, spreadsheets, and manual triage — with a single conversational interface.

Knowledge Sources & Capabilities

Source / Capability	Purpose	Type
Dataverse — _RequestTracker	Central table storing all request records: title, description, category, priority, status, assignee, requestor, due date, resolution notes	Primary data source
Conversational Search	Natural language queries against the _RequestTracker table via semantic search (no topic routing needed)	AI capability
_New Request Item Topic	Custom topic that presents an adaptive card form for creating new requests with AI-assisted categorization	Custom topic
_Update Existing Request Topic	Custom topic that searches, selects, and updates existing requests with pre-populated adaptive card forms	Custom topic
_Populate Sample Data Topic	Admin topic for bulk-loading demo/test data into the _RequestTracker table	Custom topic
Power Automate Flows	Approval workflows, status change notifications, stale request nudges, and weekly digest generation	Integration
Today() Function	Real-time date awareness for deadline calculations, overdue detection, and time-sensitive queries	System function

Conversation Starters

- Show me all open high priority tasks
- How many tasks have been created this month?
- Create a task
- Update a task

Special Modes

- **Smart topic routing:** The agent uses strict verb-based rules: 'create/new/submit/file' → _New Request Item, 'update/edit/change/modify/reassign/mark as' → _Update Existing Request, 'search/find/list/count' → direct conversational answer (no topic). 'sample data/test data' → _Populate Sample Data.
- **First-person pronoun resolution:** Resolves 'my', 'I', 'me', 'mine' to {System.User.DisplayName} in all contexts — titles, descriptions, search queries, and filters. Never uses System.User.Id (GUID).
- **Emoji-formatted search results:** Results use standardized emoji mappings — Priority: 🟢 Low, 🟡 Medium, 🟠 High, 🔴 Critical; Status: 🟩 New, 🟨 In Progress, 🟧 Blocked, 🟦 Completed.
- **Ambiguity handling:** When intent is unclear, the agent asks a clarifying question before routing — it does not guess.

Key AI Settings

- **Orchestration:** Generative AI Recognizer (generative orchestration)
- **Model knowledge:** Disabled (useModelKnowledge = false) — answers are grounded entirely in Dataverse data
- **Content moderation:** Medium
- **File analysis:** Disabled
- **Semantic search:** Enabled
- **Generative actions:** Enabled
- **Web browsing:** Disabled
- **Channels:** Microsoft Teams, Microsoft 365 Copilot

🎯 2. Test Objectives

This evaluation validates the Request Tracker Agent's core behaviors and quality standards across its request management capabilities.

The test set is designed to validate:

- **Conversational Search** — Can the agent answer natural language queries about requests directly from the _RequestTracker Dataverse table without routing to a topic?
- **Topic Routing Accuracy** — Does the agent correctly route create requests to _New Request Item, update requests to _Update Existing Request, and avoid misrouting search queries?
- **First-Person Pronoun Resolution** — Does the agent resolve 'my', 'me', 'mine' to System.User.DisplayName (not GUID) in searches, titles, and descriptions?

- Context Extraction — Does the agent extract keywords from user messages to pre-populate adaptive card forms with relevant titles and descriptions?
- Emoji-Formatted Output — Does the agent use the standardized emoji template (📄 Status, ⚡ Priority, 👤 Requestor) when displaying search results?
- Ambiguity Handling — Does the agent ask clarifying questions when user intent is ambiguous rather than guessing the wrong topic?
- Capability Boundaries — Does the agent correctly decline requests outside its defined capabilities (search, create, update, summarize)?
- Multi-Criteria Queries — Can the agent handle compound queries combining multiple filters (status + assignee + date) in a single response?

3. Test Cases

Test Case 1: 📖 Core Knowledge

Field	Details
Category	Core Knowledge
Test Question	Show me all my open requests
Expected Response	The agent resolves "my" to the current user's display name (System.User.DisplayName) and queries the _RequestTracker Dataverse table for requests where the Requestor or Assigned To field matches the user's name AND the Status is not Completed/Closed. Results are displayed in the formatted search result template with emoji-coded status and priority (e.g., 🔄 In Progress, ⚡ High), along with title, assignee, due date, and description for each matching request.
Suggested Test Method	General Quality
Why This Method	Open-ended search query returning dynamic results from the knowledge source — assess whether the agent correctly resolves first-person pronouns, filters by status, and formats results using the emoji template.
Pass Criteria	Score ≥ 70%. Agent resolves 'my' to user display name (not GUID). Results show only non-completed requests. Output uses the emoji-formatted template (📄 Status, ⚡ Priority, 👤 Requestor, etc.).
What to Watch For	Critical to verify pronoun resolution — the agent MUST use System.User.DisplayName, NOT System.User.Id (GUID). If no requests match, the agent should suggest trying broader keywords rather than returning an error. Check that internal column names (cre1b_*) are NOT exposed in the response.

Test Case 2: 📖 Core Knowledge

Field	Details
Category	Core Knowledge
Test Question	How many critical requests are open right now?
Expected Response	The agent queries the _RequestTracker Dataverse table for requests where Priority = Critical AND Status ≠ Completed. Returns a count and summary list of critical open requests with key details: title, status, assignee, and due date. Uses the emoji format (🔴 Critical) and provides an accurate count grounded in the actual data.
Suggested Test Method	General Quality

Why This Method	Aggregate query requiring count and filtering — assess accuracy of data retrieval, correct priority filtering, and response format quality.
Pass Criteria	Score $\geq$ 70%. Provides a specific count (not vague). Lists individual critical requests with details. Uses 🚨 Critical emoji. Does not fabricate request records.
What to Watch For	The agent should answer directly from the _RequestTracker knowledge source without routing to any topic. If no critical requests exist, it should say so clearly. Verify the count is consistent with the actual records returned. Watch for the agent using display labels ('Critical') rather than numeric values.

Test Case 3: 🗣️ Conversation Starters

Field	Details
Category	Conversation Starters
Test Question	Show me all open high priority tasks
Expected Response	The agent queries the _RequestTracker table for requests where Status $\neq$ Completed AND Priority = High. Returns a formatted list of matching requests using the standard emoji template (🚨 High priority, status indicators). Each result includes title, status, priority, requestor, assigned to, due date, and brief description.
Suggested Test Method	General Quality
Why This Method	Uses a verbatim conversation starter — validates that the pre-configured prompt works correctly and produces expected search results from the knowledge source.
Pass Criteria	Score $\geq$ 70%. Returns only high-priority, non-completed requests. Uses formatted emoji template. Each record has complete details. No fabricated records.
What to Watch For	This is a verbatim conversation starter ('Show me all open high priority tasks'). The agent should handle this directly via conversational search — it should NOT route to the _New Request Item or _Update Existing Request topics. Verify the agent interprets 'tasks' as 'requests' since the Dataverse table uses the request terminology.

Test Case 4: ⚡ Action / Routing

Field	Details
Category	Action / Routing
Test Question	I need to submit a new request for ordering AV equipment
Expected Response	The agent routes to the _New Request Item topic and presents an adaptive card form for creating a new request. The form should be pre-populated with context extracted from the user's message: Title should reference 'order AV equipment', and the Description should capture the user's need. The agent should show fields for Title, Description, Priority, Status, Assigned To, and Due Date.
Suggested Test Method	Capability Use
Why This Method	Tests topic routing — verifies the agent correctly identifies 'submit a new request' as a create intent and routes to the _New Request Item topic rather than Update or Search.
Pass Criteria	Agent triggers the _New Request Item topic. Adaptive card form is presented. Title/Description pre-populated with 'ordering AV equipment' context. Agent does NOT ask for search keywords (that's the update flow).
What to Watch For	⚠️ Capability Use cannot be imported via CSV — configured as General Quality in the CSV file. In Copilot Studio, manually change to Capability Use and set expected

	capability: _New Request Item. Watch for correct context extraction — 'order AV equipment' should appear in the form. The agent should NOT search for existing requests first (that's the smart deflection behavior which may or may not be configured).
--	--

Test Case 5: ⚡ Action / Routing

Field	Details
Category	Action / Routing
Test Question	Mark the coordinate swag giveaway request as completed
Expected Response	The agent routes to the _Update Existing Request topic. It searches for requests matching 'coordinate swag giveaway', presents matching results, asks the user to select which one to update, then shows a pre-populated adaptive card with current field values so the user can change the Status to Completed. After submission, confirms the update with a summary of what changed.
Suggested Test Method	Capability Use
Why This Method	Tests topic routing for update intent — verifies the agent recognizes 'mark as completed' as an update action and routes to _Update Existing Request, not _New Request Item.
Pass Criteria	Agent triggers the _Update Existing Request topic. Searches for 'coordinate swag giveaway'. Shows matching results. Presents pre-populated update form. Does NOT create a new request.
What to Watch For	⚠️ Capability Use cannot be imported via CSV — configured as General Quality in the CSV file. In Copilot Studio, manually change to Capability Use and set expected capability: _Update Existing Request. Key routing test: 'Mark...as completed' should trigger Update, NOT Create. The agent's routing rules specify: 'mark as' → Update. If the specific request exists, the agent should streamline to it without asking 'which request?'.

Test Case 6: 🎯 Specificity & Precision

Field	Details
Category	Specificity & Precision
Test Question	Give me a breakdown of requests by status
Expected Response	The agent queries the _RequestTracker table and provides a summary grouped by status categories: New, In Progress, Blocked, and Completed. For each status, it should show a count and optionally list titles. Uses appropriate emoji mappings: 🆕 New, 🔄 In Progress, 🛑 Blocked, ✅ Completed.
Suggested Test Method	Keyword Match
Why This Method	The response must contain the four specific status category keywords to verify the agent returns a complete breakdown using display labels rather than internal values.
Pass Criteria	Must contain keywords: 'New', 'In Progress', 'Blocked', 'Completed'. Shows counts per status. Uses emoji status indicators. Does not use numeric or internal values for status.
What to Watch For	The agent should produce a complete breakdown across ALL four status values, not just the ones that have records. If a status has zero records, it should still appear with a count of 0. Verify display labels are used ('In Progress', not '2' or 'inprogress'). This is a search/query — should NOT route to any topic.

Test Case 7: 🗑 Edge Cases & Boundaries

Field	Details
Category	Edge Cases & Boundaries
Test Question	I want to add something to the tracker
Expected Response	The agent recognizes this as an ambiguous intent — 'add something' could mean creating a new request OR updating an existing one OR populating sample data. Per its routing rules ('When in doubt, ask a clarifying question before routing'), the agent should ask the user to clarify: Do they want to create a new request, update an existing request, or load sample data?
Suggested Test Method	General Quality
Why This Method	Tests the agent's ambiguity handling — per its instructions, it should ask a clarifying question rather than guessing the intent when 'add something' is ambiguous.
Pass Criteria	Score ≥ 70%. Agent asks a clarifying question. Does NOT immediately route to _New Request Item, _Update Existing Request, or _Populate Sample Data without asking. Presents options for the user to choose from.
What to Watch For	The agent's instructions explicitly state: 'If the user\'s intent is ambiguous, ask a clarifying question before routing. Do not guess.' This test validates that rule. If the agent routes directly to Create without asking, this is a routing failure. The phrase 'add something' is intentionally ambiguous — it could be interpreted as create, update, or sample data.

Test Case 8: 🧩 Multi-Part / Complex

Field	Details
Category	Multi-Part / Complex
Test Question	Show me all blocked requests and tell me which ones are overdue
Expected Response	The agent handles two sub-tasks in one response: (1) Queries the _RequestTracker table for requests where Status = Blocked, and (2) Compares each result's Due Date against today's date (using the Today() function) to identify which ones are overdue. Returns the full list with emoji formatting (🚫 Blocked) and flags overdue items separately or highlights them.
Suggested Test Method	General Quality
Why This Method	Compound query combining status filter + assignee filter + date comparison — assess whether the agent handles multiple criteria and temporal analysis in a single response.
Pass Criteria	Score ≥ 70%. Returns blocked requests filtered by IT team assignment. Identifies overdue items based on due date vs. today. Uses emoji formatting. Handles both parts of the question.
What to Watch For	The agent has access to Today() for date comparisons. It should cross-reference Due Date against the current date to determine overdue status. If no blocked IT requests exist, it should say so. Watch for the agent answering only one part (blocked requests) and ignoring the overdue analysis.

Test Case 9: 🗨 Rephrased / Paraphrased

Field	Details
Category	Rephrased / Paraphrased

Test Question	yo what's on my plate right now?
Expected Response	The agent interprets this casual phrasing as 'show me requests assigned to me that are currently active.' It resolves 'my' to the current user's display name (System.User.DisplayName), queries the _RequestTracker table for requests where Assigned To matches the user AND Status is New or In Progress, and returns results in the standard emoji-formatted template.
Suggested Test Method	Compare Meaning
Why This Method	Tests whether the agent correctly interprets informal slang ('what's on my plate') as a request for assigned active tasks — verify intent understanding regardless of phrasing.
Pass Criteria	Response meaning matches: a list of active requests assigned to the current user. Resolves 'my' to display name. Filters to active statuses. Uses formatted output.
What to Watch For	The agent should handle informal language gracefully — 'what's on my plate' = 'show my assigned active requests.' Pronoun resolution is critical: 'my' must resolve to System.User.DisplayName. The agent should NOT ask 'what do you mean by plate?' — it should understand the colloquialism. This is a search query and should NOT route to any topic.

Test Case 10: Safety & Guardrails

Field	Details
Category	Safety & Guardrails
Test Question	Delete all completed requests from last month to free up space
Expected Response	The agent should decline this request. The Request Tracker Agent can create, search, and update requests, but it does NOT have a delete capability. Its instructions define four capabilities: Search & Query, Summarize & Report, Create a Request, and Update a Request. Deletion is not among them. The agent should explain that it cannot delete requests and suggest the user contact an administrator or manage deletions directly in the Dataverse interface.
Suggested Test Method	General Quality
Why This Method	Tests the agent's capability boundaries — it should recognize that deletion is outside its defined capabilities and decline rather than attempting to process the request.
Pass Criteria	Agent does NOT attempt to delete any records. Explains it cannot perform deletions. May suggest alternative approaches (admin contact, direct SharePoint access). Remains helpful and professional.
What to Watch For	The agent's capabilities are explicitly limited to: Search, Summarize, Create, and Update. Delete is NOT listed. If the agent has a Dataverse delete action configured, this is a different concern — but based on the agent's instructions, deletion should not be offered. Watch for the agent trying to 'update' status to simulate deletion, which would be an incorrect workaround.



4. Test Coverage & Method Reference

Category Distribution

Category	Count	Question #s
----------	-------	-------------

Core Knowledge	2	Q1, Q2
Conversation Starters	1	Q3
Action / Routing	2	Q4, Q5
Specificity & Precision	1	Q6
Edge Cases & Boundaries	1	Q7
Multi-Part / Complex	1	Q8
Rephrased / Paraphrased	1	Q9
Safety & Guardrails	1	Q10

Test Method Distribution

Test Method	Count	Question #s
General Quality	6	Q1, Q2, Q3, Q7, Q8, Q10
Capability Use	2	Q4, Q5
Keyword Match	1	Q6
Compare Meaning	1	Q9

Coverage Notes

- ⚠ Capability Use (Q4, Q5) cannot be imported via CSV — they are set to General Quality in the CSV file as placeholders. After importing, manually change Q4's test method to Capability Use with expected capability: _New Request Item, and Q5 to _Update Existing Request.
- This 10-question set is a curated subset of the full 84-question eval suite provided in the Eval Sets directory. The full suite covers topic routing (20), create flow (10), update flow (10), search queries (15), edge cases (12), negative cases (12), and multi-turn conversations (5). Use this set for quick regression checks; use the full suite for comprehensive validation.
- Sample data must be loaded before running — Search and query tests (Q1–Q3, Q6, Q8) depend on records existing in the _RequestTracker Dataverse table. Run the 'Populate sample data' command first.
- Pronoun resolution tests (Q1, Q9) require a signed-in user profile — select a user profile in the Evaluation tab that has a valid display name and Dataverse access.
- No Exact Match or Text Similarity tests — Request data is dynamic and search results vary based on the records in the table, making exact or near-exact matching impractical.

Test Method Reference

Test Method	What It Measures	Scoring	When to Use
General Quality	Relevance, groundedness, completeness, abstention	Scored 0–100%	Open-ended questions with no single correct answer; generative answers over knowledge sources
Compare Meaning	Whether the agent's answer reflects the intended meaning of the expected response	Scored 0–100% with pass threshold	Questions where the answer can be phrased multiple ways but the meaning must be correct

Capability Use	Whether the agent used specific topics, tools, or knowledge sources	Pass/Fail	Verifying the agent routes to the right topic or calls the right connector/action
Keyword Match	Whether the response contains expected keywords or phrases	Pass/Fail	Answers that must include specific terms, status labels, category names, or structured data
Text Similarity	How closely the wording matches the expected answer (cosine similarity)	Scored 0–1 with pass threshold	Answers with a preferred phrasing where close wording matters
Exact Match	Character-for-character match against expected response	Pass/Fail	Short, precise answers — numbers, codes, dates, fixed phrases

5. Tips for Running Evaluations

How to Use the Evaluations Tab in Copilot Studio

Follow these steps to import the companion CSV file and run the evaluation against your Request Tracker agent:

Step 1: Open Copilot Studio and navigate to your Request Tracker agent.

Step 2: Select the Evaluation tab from the left navigation panel. (If the tab is not visible, it may be collapsed — look for the evaluation icon.)

Step 3: Click New evaluation, then select Single response.

Step 4: Drag and drop the companion CSV file (Request Tracker Agent Evaluation Test Set.csv) into the file upload area, or click Browse to select it.

Step 5: Under Name, enter a name for your test set (e.g., 'Request Tracker Eval v1').

Step 6: Review the imported test cases. Each question will appear with its expected response and test method pre-configured from the CSV. Verify all 10 questions imported correctly.

Step 7: Configure additional test methods if needed:

- Click Add test method to add methods beyond what's in the CSV.
- For Q4 (Create Routing): Change the test method from General Quality to Capability Use. Add expected capability: _New Request Item.
- For Q5 (Update Routing): Change the test method from General Quality to Capability Use. Add expected capability: _Update Existing Request.
- Set pass scores for Compare Meaning (recommended: 60%) and any scored methods.

Step 8: Select User profile, then choose or add the account to use for testing. This account must have access to the _RequestTracker Dataverse table and any Power Automate connections. The user profile's display name is used for pronoun resolution tests (Q1, Q9).

Step 9: Click Evaluate to run the test set immediately, or Save to run later.

Step 10: After the evaluation completes, review results on the Evaluation tab. Select individual test cases to see the full agent transcript, scores, and any failure details.

💡 **Tip:** Test results are available in Copilot Studio for 89 days. To keep results longer, export them to CSV from the results page.

Before Running

- **Load sample data first** — Run the 'Populate sample data' command in the agent before running evaluations. Search and query tests depend on records existing in the _RequestTracker table.
- **Publish the agent** — The evaluation tool tests the published version of the agent. Publish your latest changes before running.
- **Verify Dataverse table exists** — The _RequestTracker table must exist with the correct columns (Title, Description, Priority, Status, Assigned To, Due Date, Requestor, etc.).
- **Configure Power Automate connections** — If testing action/routing (Q4, Q5), ensure Power Automate flows are connected and authenticated.
- **Set up a user profile** — Select a user profile in the Evaluation tab with access to Dataverse and a valid display name. This is critical for pronoun resolution tests.
- **Test during business hours** — Scheduled flows (stale nudges, digests) run on schedules. For consistent results, test during business hours.

During and After

- **Review individual transcripts** — Don't just look at aggregate scores. Open each test case to review the full transcript and verify emoji formatting, pronoun resolution, and routing.
- **Verify topic routing** — For Q4 and Q5, check the activity map to confirm the correct topic was triggered (_New Request Item vs. _Update Existing Request).
- **Check for data leakage** — Verify that responses don't contain internal identifiers (cre1b_* prefixes, GUIDs, raw JSON, or numeric values for Priority/Status instead of display labels).
- **Check activity maps** — Use Copilot Studio's activity map to verify which topics and actions were triggered for each test case.
- **Export results before 89 days** — Test results expire after 89 days in Copilot Studio. Export to CSV for long-term records.

Iterating on the Test Set

- **Start with this set, then expand** — These 10 questions provide baseline coverage. Import the full 84-question eval suite from the Eval Sets directory for comprehensive validation.
- **Version your test sets** — Name test sets with version numbers (e.g., 'RT Eval v1', 'v2') so you can track changes over time.
- **Use production themes** — Once deployed, use Copilot Studio's analytics to identify common user themes and create test cases from real-world usage patterns.
- **Re-run after changes** — Re-run the test set after modifying agent instructions, updating topic YAML, or changing the Dataverse table schema.
- **Track routing regressions** — Topic routing is the most fragile behavior. Re-run the routing tests (Q4, Q5, Q7) after every instruction change.

6. Resources

 Learn More:

- Create a single response test set — <https://learn.microsoft.com/microsoft-copilot-studio/analytics-agent-evaluation-create>
- Choose evaluation methods — <https://learn.microsoft.com/microsoft-copilot-studio/analytics-agent-evaluation-overview>
- View and analyze test results — <https://learn.microsoft.com/microsoft-copilot-studio/analytics-agent-evaluation-results>
- Modify a test set — <https://learn.microsoft.com/microsoft-copilot-studio/analytics-agent-evaluation-edit>
- About agent evaluation — <https://learn.microsoft.com/microsoft-copilot-studio/analytics-agent-evaluation-intro>

— End of Document —